

Predictions and Visualization

1. Overview

To evaluate the real-world effectiveness of the XGBoost model in forecasting hourly car-sharing demand, we performed visual comparisons between the model's predictions and actual observed values. These visualizations help validate the model's accuracy, generalization ability, and behavior under varying demand patterns.

2. Performance Metrics

Metric **Value**

RMSE 167.73


MAE 98.44

MAPE (%) 38.65

R² Score 0.933

- The XGBoost model explains **93.3%** of the variance in the data.
- Error metrics (RMSE and MAE) are significantly lower than the target variable's standard deviation (~645), indicating high prediction accuracy.

3. Actual vs Predicted Plot



- **Description:**
This line plot compares actual car rental counts against the model's predictions for the first 100 time intervals.
- **Observations:**

- The predicted line closely follows the actual trend, capturing key fluctuations, including morning and evening peaks.
 - Minor deviations are observed during extreme demand changes, but the overall alignment is strong.
 - **Interpretation:**
The close match demonstrates that the model generalizes well and reliably captures time-based demand patterns.
-

4. Final Output Table

	Actual	Predicted	Accuracy
6056	1728	1960.544678	0.865426
5556	822	768.247620	0.934608
5990	658	675.289795	0.973724
7674	2716	2484.060791	0.914603
3319	1083	734.010437	0.677757

Full prediction output has been saved and included as a .csv file for verification and future deployment.

Conclusion

The prediction visualizations and residual analysis confirm that the XGBoost model performs exceptionally well for hourly car-sharing demand forecasting. The alignment between actual and predicted values, combined with balanced error distribution and interpretable feature importance, ensures its reliability for real-world application.